

# MIX10: A MATLAB TO X10 COMPILER

*by*

*Vineet Kumar*

School of Computer Science  
McGill University, Montréal

Thursday, October 31st 2013

A THESIS SUBMITTED TO THE FACULTY OF GRADUATE STUDIES AND RESEARCH  
IN PARTIAL FULFILLMENT OF THE REQUIREMENTS FOR THE DEGREE OF  
MASTER OF SCIENCE

Copyright © 2013 Vineet Kumar



# **Abstract**

TBD



# Résumé

TBD



# Acknowledgements

TBD





# Table of Contents

<b>Abstract</b>	<b>i</b>
<b>Résumé</b>	<b>iii</b>
<b>Acknowledgements</b>	<b>v</b>
<b>Table of Contents</b>	<b>vii</b>
<b>List of Figures</b>	<b>ix</b>
<b>List of Tables</b>	<b>xi</b>
<b>1 Introduction</b>	<b>1</b>
1.1 Contributions . . . . .	2
1.2 Thesis Outline . . . . .	4
<b>2 Introduction to X10 programming language</b>	<b>7</b>
2.1 Overview of X10's key sequential features . . . . .	7
2.1.1 Object-oriented features . . . . .	8
2.1.2 Statements . . . . .	8
2.1.3 Types . . . . .	8
2.2 Overview of X10's concurrency features . . . . .	8
2.2.1 The APGAS Model . . . . .	8
2.2.2 Async Construct . . . . .	8
2.2.3 Finish Construct . . . . .	8

2.2.4	Atomic and When Constructs . . . . .	8
2.2.5	Places and At Construct . . . . .	8
2.3	Overview of X10's implementation and runtime . . . . .	8
<b>3</b>	<b>Background and High level design</b>	<b>9</b>
<b>4</b>	<b>Static analyses for performance and extended feature support</b>	<b>11</b>
<b>5</b>	<b>Code generation</b>	<b>13</b>
<b>6</b>	<b>Techniques for efficient compilation of MATLAB arrays</b>	<b>15</b>
6.1	The Function Collection Object . . . . .	16
6.2	The Interprocedural Analysis Framework . . . . .	17
6.2.1	Contexts . . . . .	18
6.2.2	Call Strings . . . . .	19
6.2.3	Callsite . . . . .	20
6.2.4	Recursion . . . . .	21
6.3	Summary . . . . .	23
<b>7</b>	<b>Evaluation</b>	<b>25</b>
<b>8</b>	<b>Related work</b>	<b>27</b>
<b>9</b>	<b>Conclusions and Future Work</b>	<b>29</b>
 <b>Appendices</b>		
<b>A</b>	<b>XML structure for builtin framework</b>	<b>31</b>
<b>B</b>	<b>isComplex analysis Propagation Language</b>	<b>33</b>
<b>C</b>	<b>MIX10 IR Grammar</b>	<b>35</b>
	<b>Bibliography</b>	<b>37</b>

## List of Figures

6.1	A small program where <i>main</i> calls <i>f</i> calls <i>g</i> . . . . .	19
6.2	A small program showing two calls to a function . . . . .	20
6.3	Multiple possible callsites from one statement . . . . .	21
6.4	A recursive example . . . . .	22
6.5	Example program showing an infinite chain of calls. . . . .	23



# List of Tables

6.1	The different kinds of Function Collection objects. . . . .	17
-----	---	----



# Chapter 1

## Introduction

---

MATLAB is a popular numeric programming language, used by millions of scientists, engineers as well as students worldwide[Mol]. MATLAB programmers appreciate the high-level matrix operators, the fact that variables and types do not need to be declared, the large number of library and builtin functions available, and the interactive style of program development available through the IDE and the interpreter-style read-eval-print loop. However, even though MATLAB programmers appreciate all of the features that enable rapid prototyping, their applications are often quite compute intensive and time consuming. These applications could perform much more efficiently if they could be easily ported to a high performance computing system.

X10 [IBM12], on the other hand, is an object-oriented and statically-typed language which uses cilk-style arrays indexed by *Point* objects and rail-backed multidimensional arrays, and has been designed with well-defined semantics and high performance computing in mind. The X10 compiler can generate C++ or Java code and supports various communication interfaces including sockets and MPI for communication between nodes on a parallel computing system.

In this thesis we present MIX10, a source-to-source compiler that helps to bridge the gap between MATLAB, a language familiar to scientists, and X10, a language designed for high performance computing systems. MIX10 statically compiles MATLAB programs to X10 and thus allows scientists and engineers to write programs in MATLAB (or use old programs already written in MATLAB) and still get the benefits of high performance

computing without having to learn a new language. Also, systems that use MATLAB for prototyping and C++ or Java for production, can benefit from MIX10 by quickly converting MATLAB prototypes to C++ or Java programs via X10

On one hand, all the aforementioned characteristics of MATLAB make it a very user-friendly and thus popular application to develop software among a non-programmer community. On the other hand, these same characteristics make MATLAB a difficult language to compile statically. Even the de facto standard, Mathworks' implementation of MATLAB is essentially an interpreter with a *JIT accelerator* [The02] which is generally slower than statically compiled languages. GNU Octave, which is a popular open source alternative to MATLAB and is mostly compatible with MATLAB, is also implemented as an interpreter [Oct]. Lack of formal language specification, unconventional semantics and closed source make it even harder to write a compiler for MATLAB. Furthermore, the use of arrays as default data type and the dynamicity of the base types and shapes of arrays also make it harder to add support for concurrency in a static MATLAB compiler. Mathworks' proprietary solution for concurrency is the *Parallel Computing Toolbox* [Mat13], which allows users to use multicore processors, GPUs and clusters. However, this toolbox uses heavyweight worker threads and has limited scalability.

Built on top of *McLAB* static analysis framework [Doh11, DH12], MIX10, together with its set of reusable static analyses for performance optimization and extended support for MATLAB features, ultimately aims to provide MATLAB's ease of use, to benefit from the advantages of static compilation, and to expose scalable concurrency.

## 1.1 Contributions

The major contributions of this thesis are as follows:

**Identifying key challenges:** We have identified the key challenges in performing a semantics-preserving efficient translation of MATLAB to X10.

**Overall design of MIX10:** Building upon the *McLAB* frontend and analysis framework, we provide the design of the MIX10 source-to-source translator that includes a low-



level X10 IR and a template-based specialization framework for handling builtin operations.

**Static analyses:** We provide a set of reusable static analyses for performance optimization and extended support for MATLAB features. These analyses include: (1) *IntegerOkay analysis* - We provide an analysis to automatically identify variables that can be safely declared to be of type `Int` (or `Long`) without affecting the correctness of the generated X10 code. This helps to eliminate most of, otherwise necessary, typecast operations which our experiments showed to be a major performance bottleneck in the generated code; (2) *Variable renaming for type collision* - MATLAB allows a variable to hold values of different types at different points in a program. However, in statically typed languages like X10 this behaviour cannot be supported since a variable's type needs to be declared statically by the programmer and cannot be changed at any point in the program. We provide an analysis to identify and rename such variables if their different types belong to mutually exclusive UD-DU webs; and (3) *isComplex value analysis* - We designed an analysis for identification of complex numerical values in a MATLAB program. This helped us to extend MiX10 compiler to also generate X10 code for MATLAB programs that involve use of complex numerical values.

**Code generation strategies for key language constructs:** There are some very significant differences between the semantics of MATLAB and X10. A key difference is that MATLAB is dynamically-typed, whereas X10 is statically-typed. Furthermore, the type rules are quite different, which means that the generated X10 code must include the appropriate explicit type conversion rules, so as to match the MATLAB semantics. Other MATLAB features, such as multiple returns from functions, a non-standard semantics for `for` loops, and a very general range operator, must also be handled correctly. MiX10 not only supports all the key sequential constructs but also supports concurrency constructs like `parfor` and can handle vectorized instructions in a concurrent fashion. We have also designed and implemented a template-based system that allows us to generate specialized X10 code for a collection of important MATLAB builtin operations.

**Techniques for efficient compilation of MATLAB arrays:** Arrays are the core of MATLAB. All data, including scalar values are represented as arrays in MATLAB. Efficient compilation of arrays is the key for good performance. X10 provides two types of array representations for multidimensional arrays: (1) Cilk-styled, region-based arrays and (2) rail-backed *simple* arrays. We compare and contrast these two array forms for a high performance computing language in context of being used as a target language and provide techniques to compile MATLAB arrays to two different representations of arrays provided by X10.

**Working implementation and performance results:** We have implemented the MiX10 compiler over various MATLAB compiler tools provided by the *McLAB* toolkit. In the process we also implemented some enhancements to these existing tools. We provide performance results for different X10 backends over a set of benchmarks and compare them with results from other MATLAB compilers including Mathworks' MATLAB implementation and Octave.

## 1.2 Thesis Outline

This thesis is divided into 9 chapters, including this one and is structured as follows.

*Chapter 2* provides an introduction to the X10 language and describes how it compares to MATLAB from the point of view of language design. *Chapter 3* gives a description of various existing MATLAB compiler tools upon which MiX10 is implemented, presents a high-level design of MiX10, and explains the design and need of MiX10 IR. In *Chapter 4* we provide a description of the *IntegerOkay* analysis to identify variables that are safe to be declared as `Long` type, *variable renaming for type conflict* to rename variables with conflicting types in isolated UD-DU webs and *isComplex* analysis to identify complex numerical values. *Chapter 5* gives details of code generation strategies for important MATLAB constructs. In *Chapter 6* we introduce different types of arrays provided by X10, we identify pros and cons of both kinds of arrays in the context of X10 as a target language and describe code generation strategies for them. *Chapter 7* provides performance results for code generated using MiX10 for a suite of benchmarks. *Chapter 8* provides an overview

## 1.2. Thesis Outline

---

of related work and *Chapter 9* concludes and outlines possible future work.



## Chapter 2

# Introduction to X10 programming language

---

In this chapter, we describe key X10 semantics and features and contrast them with MATLAB to help readers unfamiliar with X10 and MATLAB to have a better understanding of the MIX10 compiler.

X10 is an award winning open-source programming language being developed by IBM Research. The goal of the X10 project is to provide a productive and scalable programming model for the new-age high performance computing architectures ranging from multi-core processors to clusters and supercomputers [1].

X10, like Java, is a class-based, strongly-typed, garbage-collected and object-oriented language. It uses Asynchronous Partitioned Global Space (APGAS) model to support concurrency and distribution [1]. The X10 compiler has a native backend that compiles X10 programs to C++ and a managed backend that compiles X10 programs to Java.

In contrast to X10, MATLAB is a commercially-successful, proprietary programming language that focuses on simplicity of implementing numerical computation application [1]. MATLAB is a weakly-typed, dynamic language with unconventional semantics and uses a JIT compiler backend. It provides restricted support for high performance computing via Mathworks' parallel computing toolbox [1].

## 2.1 Overview of X10's key sequential features

X10's sequential core is a container-based object-oriented language that is very similar

to that of Java or C++ []. A X10 program consists of a collection of classes, structs or interfaces, which are the top-level compilation units. Inheritance and subtyping are fairly conventional. X10 also provides very flexible arrays based on ideas in ZPL [].

### **2.1.1 Object-oriented features**

### **2.1.2 Statements**

### **2.1.3 Types**

## **2.2 Overview of X10's concurrency features**

### **2.2.1 The APGAS Model**

### **2.2.2 Async Construct**

### **2.2.3 Finish Construct**

### **2.2.4 Atomic and When Constructs**

### **2.2.5 Places and At Construct**

## **2.3 Overview of X10's implementation and runtime**

## **Chapter 3**

# **Background and High level design**

---

TBD





## **Chapter 4**

# **Static analyses for performance and extended feature support**

---



## Chapter 5

# Code generation

---

TBD



## Chapter 6

# Techniques for efficient compilation of MATLAB arrays

---

This chapter introduces the interprocedural analysis framework. We have previously introduced the builtin framework and the Tame IR. In the next chapter we will introduce the value analysis, an interprocedural analysis that uses all these tools to build a callgraph with annotated type information. In order to implement this interprocedural analysis, we have developed the interprocedural analysis framework.

The interprocedural analysis framework builds on top of the Tame IR and the MCSAF intraprocedural analysis framework. It allows the construction of interprocedural analyses by extending an intraprocedural analysis built using the MCSAF framework. This framework works together with a callgraph object implementing the correct MATLAB look up semantics. An analysis can be run on an existing callgraph object, or it can be used to build new callgraph objects, discovering new functions as the analysis runs.

In the following sections we will introduce the interprocedural analysis framework as an extension of the intraprocedural analysis framework, and how it works in tandem with callgraph objects and the lookup objects, as well as how the framework deals with recursion. To help potential analysis writers, we have indicated the names of Java classes that correspond to the contexts in bold.

## 6.1 The Function Collection Object

In order to represent callgraphs we use an object which we call **FunctionCollection**. It is, as the name suggests, a collection of nodes representing functions, indexed by so function reference objects. Objects of type **FunctionReference** act as unique identifiers for functions. They store a function's name, in which file it is contained if applicable, and what kind of function it is (primary function, subfunction, nested function, builtin function, constructor). For nested functions, it stores in which function it is contained. Function reference objects give enough information to load a function from a file.

Nodes in the function collection not only store the code of the function and a function reference; they also provide information about its environment. The node provides a MATLAB function lookup object which is able to completely resolve any function call coming from the function. It includes information about the MATLAB path environment and other functions contained in the same file. The lookup information is provided given a function name, and optionally an mclass name (to find overloaded versions); and will return a function reference allowing the loading of functions.

The lookup information allows us to build a callgraph knowing only an entry point and a path environment, and using semantics for finding functions that correspond to the way MATLAB finds functions at runtime. This is bridging the gap between a dynamic language and static compilers, which usually require specifying what source code files are required for compilation.

The simple function collection uses only the lookup information contained in its nodes to build an approximation of a callgraph, which is naturally incomplete. We have used it for the development of the Tamer framework, as it provides a simple way to generate a callgraph which excludes discovering overloaded calls and propagation of function handles. We have implemented slightly different versions of the function collection, which are described in the table below.

<b>SimpleFunctionCollection</b>	A simple callgraph object built using MATLAB lookup semantics excluding overloading. Function Handles are loaded only in the functions where the handle is created. Obviously this an incomplete callgraph, but may be used by software tools that do not need a complete callgraph, and where the simplicity can be useful.
<b>IncrementalFunctionCollection</b>	callgraph that does the same lookup as the FunctionCollection, but does not actually load functions until they are requested. This is used to build the callgraph
<b>CompleteFunctionCollection</b>	callgraph that includes call sites for every function node and correctly represents overloading can calling function handles. This is produced by the Tamer using interprocedural analyses. This callgraph can be used to build further interprocedural analysis that are not extensions of the value analysis. It can also be used as a starting point for static backends.

**Table 6.1** The different kinds of Function Collection objects.

## 6.2 The Interprocedural Analysis Framework

The interprocedural analysis framework is an extension of the intraprocedural flow analyses provided by the MCSAF framework. It is context-sensitive to aid code generation targeting static languages like FORTRAN. FORTRAN's polymorphism features are quite limited; every generated variable needs to have one specific type. The backend may thus require that every MATLAB variable has a specific known mclass at every program point. Functions may need to be specialized for different kinds of arguments, which a context-sensitive analysis provides at the analysis level.

An interprocedural analysis is a collection of interprocedural analysis nodes, called **InterproceduralAnalysisNode**, which represent a specific intraprocedural analysis for some function and some context. The context is usually a flow representation of the passed arguments. Every such interprocedural analysis node produces a result set using the contained intraprocedural analysis. An InterproceduralAnalysisNode is generic in the intraprocedural

analysis, the context and the result - these have to be defined by an actual implementation of an interprocedural analysis.

Every interprocedural analysis has an associated `FunctionCollection` object, which may initially contain only one function acting as the entry point for the program (i.e. when building a callgraph using an `IncrementalFunctionCollection`). The interprocedural analysis requires a context (argument flow set) for the entry point to the program.

### **Algorithm**

The analysis starts by creating an interprocedural analysis node for the entry point function and the associated context, which triggers the associated intraprocedural flow analysis. As the intraprocedural flow analysis encounters calls to other functions, it has to create context objects for those calls, and ask the interprocedural analysis to analyze the called functions using the given context. The call also gets added to the set of call edges associated with the interprocedural analysis node.

As the interprocedural has to analyze newly encountered calls, the associated functions are resolved, and loaded into the callgraph if necessary. The result is a complete callgraph, and an interprocedural analysis.

### **6.2.1 Contexts**

In order to implement an interprocedural analysis, one has to define a context object. These may be the flow information of the arguments of a call; but it could be any information. The analysis itself is context-sensitive, meaning that if there are multiple calls to one function with different contexts, they are all represented by different interprocedural analysis nodes. The interprocedural analysis framework never merges contexts, which would have to be done by the specific analysis if desired.

Interprocedural analysis nodes are cached. Thus if a function/context pair is called a second time, the information will be readily available.

Note that in order to completely resolve calls, the flow information and the contexts have to include `mclass` information for variables and arguments. In order to resolve calls to function handles, the contexts have to store which arguments may refer to function handles




(and which functions they refer to).

Once the complete callgraph is built, further analyses don't need to flow mclass information, because all possible calls are resolved. But this information may still be useful to obtain more accurate analysis result, by knowing which information to flow into which calls for ambiguous call sites (see Sec. 6.2.3) - that is why the value analysis presented in Chapter ?? allows extending the flow-sets, to allow flowing information for different analyses together in one analysis, and get a more precise overall result.

### 6.2.2 Call Strings


When analyzing a function  $f$  for a given context  $c_f$ , and encountering a call to some other function  $g$ , the interprocedural analysis framework suspends the analysis of  $f$  in order to analyze the encountered call. The flow analysis has to provide a context  $c_g$  for the call to  $g$ , and an intraprocedural analysis will be created that will analyze  $g$  with  $c_g$ .



**Figure 6.1** A small program where *main* calls *f* calls *g*. The call string for  $g(c_g)$  in this example may be  $main(c_{main}) : f(c_f) : g(c_g)$ .

The set of currently suspended functions (in Figure 6.3 *main* and *f*), which are awaiting results of encountered calls that need to be analyzed correspond to the callstack of these functions at runtime, at least for non-recursive programs. We call the chain of these functions, together with their contexts a **CallString**. Every function/context pair, i.e. the associated interprocedural analysis node, has an associated call string, which corresponds to one possible stack trace during runtime. Note that interprocedural analysis nodes are cached, and may be reused. Thus in the above example, if the main function also calls *g*

with context  $c_g$ , the results of the interprocedural analysis node created for the call encountered in function  $f$  will be reused.




**Figure 6.2** Here, *main* also calls *g*, also with context  $c_g$ . Since the interprocedural analysis node for  $g(c_g)$  is reused, the call string will be reused as well.

Since the interprocedural analysis node is reused, it will have the same call string. So the call string is not an exact representation of a call stack for every call, it is merely the exact representation of one possible call stack that will reach a given function/context pair. Note that for purposes of error reporting, the call string can be presented to the user as a stack trace.

### 6.2.3 Callsite

Any statement representing a call may actually represent multiple possible calls. For example a call to a function  $g$  may be overloaded, so if arguments may have different possible mclasses, different functions named  $g$  may be called. Also, because it is up to an actual analysis to define its notion of what a context is, it is possible that an analysis may decide to produce multiple contexts for one call to a function  $f$ . This would create specialized versions of a function from a single call (this is actually possible in the value analysis presented in *Chapter ??*). A third way in which a statement may represent multiple possible calls is via function handles. An `TIRArrayGet` statement may trigger a call if the represented array is actually found to be a function handle (we call the variable accessed in a `TIRArrayGet` statement an 'array' simply because it is used in an array-indexing operation, but it could be any variable). If that function handle may refer to multiple possible functions at runtime,

then the function handle access may refer to multiple possible calls.




**Figure 6.3** This figure shows examples how it is possible for one single call site to refer to multiple possible calls. This may be due to overloading, creation of multiple contexts for a single call, or function handles.

In order to be able to represent multiple possible call edges coming out of a statement, we associate any statement that includes any calls with a **Callsite** object. This callsite can store multiple possible call edges as function/context pairs, which we call a “call” in the interprocedural analysis framework. An intraprocedural analysis, in order to request the result of a call, has to request a callsite object for a calling statement. It may then request arbitrary calls from that callsite object, which will all get associated with the calling statement.

### 6.2.4 Recursion

The interprocedural analysis framework supports simple and mutual recursion by performing a fixed point iteration within the first recursive interprocedural analysis node. In order to identify recursive and mutually recursive calls we use the call strings introduced in Sec. 6.2.2. While we established that there is no guarantee which stack trace the call string represents, we know that it will always represent one possible stack trace. Since the call stacks of all recursive and mutual recursive calls must include the function, we merely need to check, for any call, whether it already exists in its call string.

If it does, we have identified a recursive call, and must perform a fixed point iteration. To do so, we label the intraprocedural analysis node associated with the recursive call (i.e. the call to  $f(c_f)$  in Figure 6.4) as recursive. This will trigger the fixed point iteration. Because we need a result for the recursive call to continue analyzing, an actual analysis implementation has to provide a default value as a first approximation, which may be just

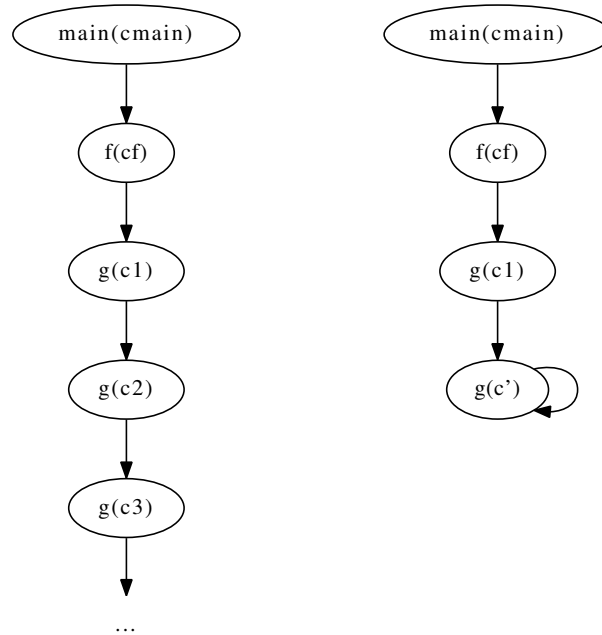


**Figure 6.4** Example of a recursive program. The call in  $s(c_s)$  to  $f(c_f)$  triggers the fixed point iteration of  $f(c_f)$ .  $f(c_f)$  is the first recursive interprocedural analysis node.

bottom. Once the intraprocedural contained in the interprocedural analysis node associated with the recursive call is completed, the result is stored as a new partial result. The analysis is then recomputed, using this new partial result for the recursive call. When a new partial result is the same as a previous partial result, we have completed the fixed point iteration. Note that the computation resulting in the new partial result uses the previous partial result for its recursive call - but since they are the same, we have made a complete analysis using the final result for the recursive call.

Note that while the fixed point iteration is being computed, all calls below the recursive call (i.e. the calls  $g(c_g)$  and  $s(c_s)$  in *Figure 6.4*) always return partial results. Thus we cannot cache the nodes and their results, and have to continuously invalidate all the corresponding interprocedural analysis nodes.

Note that the analysis treats calls to the same function with different contexts as different functions. No fixed point iteration is performed to resolve recursive calls with different contexts, because they represent different underlying intraprocedural analyses. Thus it is possible to create infinite call strings, as shown in *Figure 6.5*. It is up to the actual analysis implementation to ensure this does not happen. A simple strategy would be to ensure that there are only a finite number of possible contexts for every function. Another strategy is for the intraprocedural analysis to check the current call string before requesting a call, to



**Figure 6.5** Example of a recursive program, showing how recursive calls with different contexts can create infinite chains of calls on the left. An interprocedural analysis implementation has to catch such cases and create a finite number of contexts, as shown on the right, where the contexts  $c_2$  and onward are replaced with  $c'$ . In this case the interprocedural analysis framework will perform a fixed point iteration on  $f(c')$ .

ensure that the function to be called does not already exist in the call string. If it does, the intraprocedural analysis should push up the context to a finite representation (shown in *Figure 6.5*).

## 6.3 Summary

We have presented an interprocedural analysis framework that we hope is flexible enough to allow different kinds of full-program analyses, while powerful enough to deal with issues such as recursion and ambiguous call sites. This analysis framework is a key component of our value analysis (presented in the next chapter), and the overall callgraph construction of the Tamer.



## **Chapter 7**

### **Evaluation**

---

TBD





## **Chapter 8**

### **Related work**

---



## **Chapter 9**

# **Conclusions and Future Work**

---

TBD



## **Appendix A**

### **XML structure for builtin framework**

---

TBD



## **Appendix B**

# **isComplex analysis Propagation Language**

---

TBD





## Appendix C

### MiX10 IR Grammar

---

TBD



## Bibliography

---

- [DH12] Anton Dubrau and Laurie Hendren. Taming MATLAB. In *Proceedings of OOP-SLA 2012*, 2012, pages 503–522.
- [Doh11] Jesse Doherty. McSAF: An Extensible Static Analysis Framework for the MATLAB Language. Master’s thesis, McGill University, December 2011.
- [IBM12] IBM. X10 programming language. <http://x10-lang.org>, February 2012.
- [Mat13] MathWorks. Parallel computing toolbox, 2013. <http://www.mathworks.com/products/parallel-computing/>.
- [Mol] Cleve Moler. The Growth of MATLAB and The MathWorks over Two Decades. [http://www.mathworks.com/company/newsletters/news\\_notes/clevesco](http://www.mathworks.com/company/newsletters/news_notes/clevesco)
- [Oct] GNU Octave. <http://www.gnu.org/software/octave/index.html>.
- [The02] The Mathworks. Technology Backgrounder: Accelerating MATLAB, September 2002. <http://www.mathworks.com/company/newsletters/digest/sept02/a>